

NAME

rgbasm — language documentation

DESCRIPTION

This is the full description of the assembly language used by *rgbasm*(1). For the full description of instructions in the machine language supported by the Game Boy CPU, see *gbz80*(7).

It is advisable to have some familiarity with the Game Boy hardware before reading this document. RGBDS is specifically targeted at the Game Boy, and thus a lot of its features tie directly to its concepts. This document is not intended to be a Game Boy hardware reference.

Generally, “the linker” will refer to *gblink*(1), but any program that processes RGBDS object files (described in *rgbds*(5)) can be used in its place.

SYNTAX

The syntax is line-based, just as in any other assembler. Each line may have components in either of these orders:

- `[label:] [directive] [: comment]`
- `[label:] [instruction [:: instruction . . .]] [: comment]`

Directives are commands to the assembler itself, such as **PRINTLN**, **SECTION**, or **OPT**.

Labels tie a name to a specific location within a section (see “Labels” below). Labels are allowed before most directives, but not before **IF**, **ELIF**, **ELSE**, **ENDC**, **REPT**, **FOR**, **ENDR**, **MACRO**, or **ENDM**.

Instructions are assembled into Game Boy opcodes. Multiple instructions on one line, as well as data directives (see “Defining constant data in ROM” below), can be separated by double colons ‘: :’.

The available instructions are documented in *gbz80*(7).

Note that where an instruction requires an 8-bit register *r8*, **rgbasm** can interpret **HIGH**(*r16*) as the top 8-bit register of the given *r16*, for example, **HIGH**(**HL**) for **H**; and **LOW**(*r16*) as the bottom one, for example, **LOW**(**HL**) for **L** (except for **LOW**(**AF**), since **F** is not a valid register).

Note also that where an instruction requires a condition code *CC*, **rgbasm** can interpret **!CC** as the opposite condition code; for example, **!nz** for **Z**.

All reserved keywords (directives, instructions, registers, built-in functions, etc.) are case-insensitive; all identifiers (labels, variables, etc) are case-sensitive.

Comments are used to give humans information about the code, such as explanations. The assembler *always* ignores comments and their contents.

There are two kinds of comments, inline and block. Inline comments are anything that follows a semicolon ‘;’ not inside a string, until the end of the line. Block comments, beginning with ‘/*’ and ending with ‘*/’, can be split across multiple lines, or occur in the middle of an expression.

An example demonstrating these syntax features:

```
SECTION "My Code", ROM0 ; a directive
MyFunction:             ; a label
    push hl              ; an instruction
    /* ...and multiple instructions,
       with mixed case */
    ld a, [hli] :: LD H, [HL] :: Ld l, a
    pop /*wait for it*/ hl
    ret
```

Sometimes lines can be too long and it may be necessary to split them. To do so, put a backslash at the end of the line:

```
DB 1, 2, 3, \
   4, 5, 6, \ ; Put it before any comments
```

```

          7, 8, 9
DB "Hello, \ ; Space before the \ is included
world!"      ; Any leading space is included

```

Symbol interpolation

Symbols with string or numeric values can be “interpolated” by writing them inside `{braces}`. This will paste the symbol’s contents as if they were part of the source file. If it is a string symbol, its characters are simply inserted as-is. If it is a numeric symbol, its value is converted to hexadecimal notation with a dollar sign ‘\$’ prepended.

Symbol interpolations can be nested, too.

```

DEF topic EQU "life, the universe, and \"everything\""
DEF meaning EQU "answer"
; Defines answer = 42
DEF {meaning} = 42
; Prints "The answer to life, the universe, and "everything" is $2A"
PRINTLN "The {meaning} to {topic} is {{meaning}}"
PURGE topic, meaning, {meaning}

```

Symbols can be interpolated even in contexts that disable automatic expansion of string constants: that is, name will be expanded in all of `DEF({name})`, `DEF {name} EQU/=EQU/ETC . . .`, `REDEF {name} EQU/=EQU/ETC . . .`, `FOR {name}`, . . ., `PURGE {name}`, and `MACRO {name}`, even though it won’t be in `DEF(name)`, `PURGE {name}`, etc.

It’s possible to change the way symbols are printed by specifying a print format like so: `{fmt:symbol}`. The `fmt` specifier consists of parts, which must be in the following order: `<sign><exact><align><pad><width><frac><prec><type>`. All the parts are optional except the required `<type>`. These parts are:

Part Meaning

`<sign>` May be ‘+’ or ‘-’. If specified, prints this character in front of non-negative numbers.

`<exact>` May be ‘#’ (only allowed for non-decimal types). If specified, prints the value in an “exact” format: with a base prefix (‘\$’, ‘&’, or ‘%’) for non-decimal integer types (‘x’ / ‘X’, ‘o’, or ‘b’); with a ‘q’ precision suffix for fixed-point numbers; or with ‘\’ escape characters (but no enclosing quotes) for strings.

`<align>` May be ‘-’. If specified, aligns left instead of right.

`<pad>` May be ‘0’. If specified, pads right-aligned numbers with zeros instead of spaces.

`<width>` May be one or more ‘0’ – ‘9’. If specified, pads the value to this width, right-aligned with spaces by default.

`<frac>` May be ‘.’ followed by zero or more ‘0’ – ‘9’. If specified, prints this many fractional digits of a fixed-point number. Defaults to 5 digits, maximum 255 digits. (A ‘.’ followed by zero ‘0’ – ‘9’ prints zero fractional digits and no decimal point.)

`<prec>` May be ‘q’ followed by one or more ‘0’ – ‘9’. If specified, prints a fixed-point number at this precision. Defaults to the current `-Q` option.

`<type>` Specifies the type of value.

Valid types are:

Type	Format	Example
‘d’	Signed decimal	-42’
‘u’	Unsigned decimal	4294967254’
‘x’	Lowercase hexadecimal	2a’
‘X’	Uppercase hexadecimal	2A’
‘b’	Binary	101010’
‘o’	Octal	52’

```
'f    Fixed-point      1234.56789'
's    String           string contents'
```

Examples:

```
SECTION "Test", ROM0[2]
X:                ; This works with labels**whose address is known**
DEF Y = 3         ; This also works with variables
DEF SUM EQU X + Y ; And likewise with numeric constants
; Prints "%0010 + $3 == 5"
PRINTLN "{#05b:X} + {#x:Y} == {d:SUM}"

rsset 32
DEF PERCENT rb 1  ; Same with offset constants
DEF VALUE = 20
DEF RESULT = MUL(20.0, 0.32)
; Prints "32% of 20 = 6.40"
PRINTLN "{d:PERCENT}% of {d:VALUE} = {f:RESULT}"

DEF WHO EQU$ STRLWR("WORLD")
; Prints "Hello world!"
PRINTLN "Hello {s:WHO}!"
```

Although, for these examples, **STRFMT** would be more appropriate; see “String expressions” below.

EXPRESSIONS

There are two types of expressions: numeric and string.

Numeric expressions are always evaluated using signed 32-bit math. In Boolean logic contexts, zero is considered to be the only “false” number, and all non-zero numbers (including negative) are “true”.

An expression is said to be “constant” if **rgbasm** knows its value. This is generally always the case, unless a label is involved, as explained in the “SYMBOLS” section. However, some operators can be constant even with non-constant operands, as explained in “Operators” below.

Directives generally require constant expressions: for example, **REPT** requires the number of repetitions to be known at assembly time.

Numeric literals

rgbasm supports a variety of numeric literals.

Format type	Prefixes	Accepted characters
Decimal	none	0123456789
Hexadecimal	\$, 0x, 0X	0123456789ABCDEF
Octal	&, 0o, 0O	01234567
Binary	%, 0b, 0B	01
Fixed-point	none	01234.56789
Precise fixed-point	none	12.34q8
Character constant	none	'A'
Game Boy graphics	`	0123

Underscores are also accepted in numbers, except at the beginning of one. This can be useful for grouping digits, like 123_456 or %1100_1001.

The “character constant” form yields the value the character maps to in the current charmap. For example, by default (refer to *ascii(7)*) ‘A’ yields 65. A character constant must represent a single value, so it cannot include multiple characters, or characters which map to multiple values. See “Character maps” for information on charmaps, and “String expressions” for information on escape characters allowed in character constants.

The last one, Game Boy graphics, expects up to eight digits between 0 and 3, corresponding to pixels' two-bit shade values. The resulting numeric value is the two bytes of tile data which would produce that row of pixels. For example, '01012323' is equivalent to '\$0F55'.

In place of a numeric literal, you can also use a numeric symbol's name, which is implicitly replaced with its value.

Operators

You can use these operators in numeric expressions (listed from highest to lowest precedence):

Operator	Meaning
()	Grouping
FUNC()	Built-in function call
**	Exponentiation
+ - ~ !	Unary plus, unary minus (negation), complement (bitwise negation), and Boolean negation
* / %	Multiplication, division (rounding down), and modulo (remainder)
<< >> >>>	Bit shifts (left, sign-extended right, zero-extended right)
& ^	Bitwise AND/OR/XOR
+ -	Addition and subtraction
== != < > <= >=	Comparisons
&&	Boolean AND
	Boolean OR

'**' raises a number to a non-negative power. It is the only *right-associative* operator, meaning that $p ** q ** r$ is equal to $p ** (q ** r)$, not $(p ** q) ** r$. All other binary operators are left-associative.

'~' complements a value by inverting all 32 of its bits.

'%' is used to get the remainder of the corresponding division, so that $x / y * y + x \% y == x$ is always true. The result has the same sign as the divisor. This makes $x \% y$ equal to $(x + y) \% y$ or $(x - y) \% y$.

Shifting works by shifting all bits in the left operand either left ('<<') or right ('>>') by the right operand's amount. When shifting left, all newly-inserted bits are reset; when shifting right, they are copies of the original most significant bit instead. This makes 'a << b' and 'a >> b' equivalent to multiplying and dividing by 2 to the power of b, respectively.

Comparison operators return 0 if the comparison is false, and 1 otherwise.

Unlike in many other languages, and for technical reasons, **rgbasm** still evaluates both operands of '&&' and '||'.

The operators '&&' and '&' with a zero constant as either operand will be constant 0, and '||' with a non-zero constant as either operand will be constant 1, even if the other operand is non-constant.

'!' returns 1 if the operand was 0, and 0 otherwise. Even a non-constant operand with any non-zero bits will return 0.

Integer functions

Besides operators, there are also some functions which have more specialized uses:

Name	Operation
HIGH (n)	Equivalent to '(n & \$FF00) >> 8'.
LOW (n)	Equivalent to 'n & \$FF'.
BITWIDTH (n)	Returns the number of bits necessary to represent n. Some useful formulas: BITWIDTH (n) - 1 equals $\lceil \log_{\text{sub } 2} (n) \rceil$; BITWIDTH (n - 1) equals $\lceil \log_{\text{sub } 2} (n) \rceil$; and $32 - \text{BITWIDTH}(n)$ equals $\text{roman_clz}(n)$, the count of leading zero bits in the binary representation of n.

TZCOUNT(*n*) Returns $\$roman\ ctz\ (n)\ \$$, the count of trailing zero bits in the binary representation of *n*.
 delim off

Fixed-point expressions

Fixed-point numbers are technically just integers, but conceptually they have a decimal point at a fixed location (hence the name). This gives them increased precision, at the cost of a smaller range, while remaining far cheaper to manipulate than floating-point numbers (which **rgbasm** does not support).

The default precision of all fixed-point numbers is 16 bits, meaning the lower 16 bits are used for the fractional part; so they count in 65536ths of 1.0. This precision can be changed with the **-Q** command-line option, and/or by **OPT Q** (see “Changing options while assembling”). An individual fixed-point literal can specify its own precision, overriding the current default, by appending a “q” followed by the number of fractional bits: for example, **789.25q8** is equal to $\$000315_40\ \$$ ($\$ = 789.25 * 2^{sup\ 8}$).

Since fixed-point values are still just integers, you can use them in normal integer expressions. You can easily truncate a fixed-point number into an integer by shifting it right by the number of fractional bits, or by dividing it by 1.0. It follows that you can convert an integer to a fixed-point number by shifting it left that same amount, or by multiplying it by 1.0. For example, $123.0 / 1.0 == 123$, and $123 * 1.0 == 123.0$.

Note that the current number of fractional bits can be computed as **TZCOUNT**(1.0).

The following functions are designed to operate with fixed-point numbers (which must be known constant):

Name	Operation
DIV (<i>x</i> , <i>y</i>)	Fixed-point division
MUL (<i>x</i> , <i>y</i>)	Fixed-point multiplication
FMOD (<i>x</i> , <i>y</i>)	Fixed-point modulo
POW (<i>x</i> , <i>y</i>)	$\$x\ sup\ y\$$
LOG (<i>x</i> , <i>y</i>)	Logarithm of $\$x\$$ to the base $\$y\$$
ROUND (<i>x</i>)	Round $\$x\$$ half away from zero to the nearest integer
CEIL (<i>x</i>)	Round $\$x\$$ up to the nearest integer
FLOOR (<i>x</i>)	Round $\$x\$$ down to the nearest integer
SIN (<i>x</i>)	Sine of $\$x\$$
COS (<i>x</i>)	Cosine of $\$x\$$
TAN (<i>x</i>)	Tangent of $\$x\$$
ASIN (<i>x</i>)	Inverse sine of $\$x\$$
ACOS (<i>x</i>)	Inverse cosine of $\$x\$$
ATAN (<i>x</i>)	Inverse tangent of $\$x\$$
ATAN2 (<i>y</i> , <i>x</i>)	Angle between $\$(x, y)\ \$$ and $\$(1, 0)\ \$$

delim off

There are no functions for fixed-point addition and subtraction, because the ‘+’ and ‘-’ operators can add and subtract pairs of fixed-point operands.

Note that some operators or functions are meaningful when combining integers and fixed-point values. For example, $2.0 * 3$ is equivalent to **MUL**(**2.0**, **3.0**), and $6.0 / 2$ is equivalent to **DIV**(**6.0**, **2.0**). Be careful and think about what the operations mean when doing this sort of thing.

All of these fixed-point functions can take an optional final argument, which is the precision to use for that one operation. For example, **MUL**(**6.0q8**, **7.0q8**, **8**) will evaluate to **42.0q8** no matter what value is set as the current **Q** option. **rgbasm** does not check precisions for consistency, so nonsensical input like **MUL**(**4.2q8**, **6.9q12**, **16**) will produce a nonsensical (but technically correct) result: “garbage in, garbage out”.

The **FMOD** function is used to get the remainder of the corresponding fixed-point division. The result has the same sign as the *dividend*; this is the opposite of how the integer modulo operator ‘%’ works!

The trigonometry functions (**SIN**, **COS**, **TAN**, etc) are defined in terms of a circle divided into 1.0 “turns” delim \$\$ (equal to \$2 pi\$ radians, or 360 degrees). delim off

These functions are useful for automatic generation of various tables. For example:

```
; Generate a table of 128 sine values
; from sin(0.0) included to sin(0.5) excluded,
; with amplitude scaled from [-1.0, 1.0] to [0.0, 128.0],
; then divided by 1.0 to round down to integer values.
FOR angle, 0.0, 0.5, 0.5 / 128
    db MUL(SIN(angle) + 1.0, 128.0 / 2) / 1.0
ENDR
```

String expressions

The most basic string expression is a string literal: any number of characters contained in double quotes ("for instance"). The backslash character ‘\’ is special in that it causes the character following it to be “escaped”, meaning that it is treated differently from normal. There are a number of escape sequences you can use within a string:

Sequence	Meaning
‘\\	Backslash (escapes the escape character itself)’
‘\"	Double quote (does not terminate a string)’
‘\’	Single quote (does not terminate a character literal)’
‘\{	Open curly brace (does not start interpolation)’
‘\}	Close curly brace (does not end interpolation)’
‘\n	Newline (ASCII \$0A)’
‘\r	Carriage return (ASCII \$0D)’
‘\t	Tab (ASCII \$09)’
‘\0	Null (ASCII \$00)’

Multi-line string literals are contained in triple quotes ("\"for instance\""). Escape sequences work the same way in multi-line strings; however, literal newline characters will be included as-is, without needing to escape them with ‘\r’ or ‘\n’.

Raw string literals are prefixed by a hash ‘#’. Inside them, backslashes and braces are treated like regular characters, so they will not be expanded as macro arguments, interpolated symbols, or escape sequences. For example, the raw string #"\t\1{s}\\" is equivalent to the regular string "\\t\\1\\{s}\\\". (Note that this prevents raw strings from including the double quote character.) Raw strings also may be contained in triple quotes for them to be multi-line, so they can include literal newline or quote characters (although still not three quotes in a row).

You can use the ‘++’ operator to concatenate two strings. "str" ++ "ing" is equivalent to "string", or to STRCAT("str", "ing").

You can use the ‘===’ and ‘!==' operators to compare two strings. "str" === "ing" is equivalent to STRCMP("str", "ing") == 0, and "str" !== "ing" is equivalent to STRCMP("str", "ing") != 0.

The following functions operate on string expressions, and return strings themselves:

Name	Operation
STRCAT (<i>strs...</i>)	Concatenates <i>strs</i> .
STRUPR (<i>str</i>)	Returns <i>str</i> with all ASCII letters (a-z) in uppercase.
STRLWR (<i>str</i>)	Returns <i>str</i> with all ASCII letters (A-Z) in lowercase.
STRSLICE (<i>str</i> , <i>start</i> , <i>stop</i>)	Returns a substring of <i>str</i> starting at <i>start</i> and ending at <i>stop</i> (exclusive). If <i>stop</i> is not specified, the substring continues to the end of <i>str</i> .

STRRPL(*str*, *old*, *new*) Returns *str* with each occurrence of the substring *old* replaced with *new*.
STRFMT(*fmt*, *args*...) Returns the string *fmt* with each %spec pattern replaced by interpolating the format *spec* (using the same syntax as “Symbol interpolation”) with its corresponding argument in *args* (‘%%’ is replaced by the ‘%’ character).
STRCHAR(*str*, *idx*) Returns the substring of *str* for the charmap entry at *idx* with the current charmap. (*idx* counts charmap entries, not characters.)

The following functions take varying operands, and return strings:

REVCHAR(*vals*...) Returns the string that is mapped to *vals* with the current charmap. If there is no unique charmap entry for *vals*, an error occurs.
READFILE(*name*, *max*) Returns the contents of the file *name* as a string. Reads up to *max* bytes, or the entire contents if *max* is not specified. If the file isn’t found in the current directory, the include-path list passed to *rgbasm*(1)’s **-I** option on the command line will be searched.

The following functions operate on string expressions, but return integers:

Name	Operation
STRLEN (<i>str</i>)	Returns the number of characters in <i>str</i> .
STRCMP (<i>str1</i> , <i>str2</i>)	Compares <i>str1</i> and <i>str2</i> according to ASCII ordering of their characters. Returns -1 if <i>str1</i> is lower than <i>str2</i> , 1 if <i>str1</i> is greater than <i>str2</i> , or 0 if they match.
STRFIND (<i>str</i> , <i>sub</i>)	Returns the first index of <i>sub</i> in <i>str</i> , or -1 if it’s not present.
STRRFIND (<i>str</i> , <i>sub</i>)	Returns the last index of <i>sub</i> in <i>str</i> , or -1 if it’s not present.
BYTELEN (<i>str</i>)	Returns the number of bytes in <i>str</i> . (Non-ASCII characters can be multiple bytes.)
STRBYTE (<i>str</i> , <i>idx</i>)	Returns the byte value at <i>idx</i> in <i>str</i> .
INCHARMAP (<i>str</i>)	Returns 1 if <i>str</i> has an entry in the current charmap, or 0 otherwise.
CHARLEN (<i>str</i>)	Returns the number of charmap entries in <i>str</i> with the current charmap.
CHARCMP (<i>str1</i> , <i>str2</i>)	Compares <i>str1</i> and <i>str2</i> according to their charmap entry values with the current charmap. Returns -1 if <i>str1</i> is lower than <i>str2</i> , 1 if <i>str1</i> is greater than <i>str2</i> , or 0 if they match.
CHARSIZE (<i>char</i>)	Returns how many values are in the charmap entry for <i>char</i> with the current charmap.
CHARVAL (<i>char</i> , <i>idx</i>)	Returns the value at <i>idx</i> of the charmap entry for <i>char</i> . If <i>idx</i> is not specified, <i>char</i> must have a single value, which is returned.

Note that indexes count starting from 0 at the beginning, or from -1 at the end. The characters of a string are counted by **STRLEN**; the charmap entries of a string are counted by **CHARLEN**; and the values of a charmap entry are counted by **CHARSIZE**.

Character maps

When writing text strings that are meant to be displayed on the Game Boy, the character encoding in the ROM may need to be different than the source file encoding. For example, the tiles used for uppercase letters may be placed starting at tile index 128, which differs from ASCII starting at 65.

Character maps allow mapping strings or character literals to arbitrary sequences of numbers:

```
CHARMAP "A", 42
CHARMAP ':)', 39
CHARMAP "<br>", 13, 10
CHARMAP '&euro;', $20ac
```

This would result in db "Amen :)
" being equivalent to db 42, 109, 101, 110, 32, 39, 13, 10, and dw "25€" being equivalent to dw 50, 53, \$20ac.

Character mappings are matched greedily, so the longest applicable one will be mapped in a string. Any characters in the string without defined mappings will be copied directly, using the source file’s encoding of characters to bytes.

It is possible to create multiple character maps and then switch between them as desired. This can be used to encode debug information in ASCII and use a different encoding for other purposes, for example. Initially, there is one character map called ‘main’ and it is automatically selected as the current character map from the beginning. There is also a character map stack that can be used to save and restore which character map is currently active.

Command	Meaning
NEWCHARMAP <i>name</i>	Creates a new, empty character map called <i>name</i> and switches to it.
NEWCHARMAP <i>name, basename</i>	Creates a new character map called <i>name</i> , copied from character map <i>basename</i> , and switches to it.
SETCHARMAP <i>name</i>	Switch to character map <i>name</i> .
PUSHC	Push the current character map onto the stack.
PUSHC <i>name</i>	Push the current character map onto the stack and switch to character map <i>name</i> .
POPC	Pop a character map off the stack and switch to it.

Note: Modifications to a character map take effect immediately from that point onward.

Other functions

There are a few other functions that do things beyond numeric or string operations:

Name	Operation
DEF (<i>symbol</i>)	Returns 1 if <i>symbol</i> has been defined, 0 otherwise. String constants are not expanded within the parentheses.
ISCONST (<i>arg</i>)	Returns 1 if <i>arg</i> ’s value is known by RGBASM (e.g. if it can be an argument to IF), or 0 if only RGLINK can compute its value.
BANK (<i>arg</i>)	Returns a bank number. If <i>arg</i> is the symbol @, this function returns the bank of the current section. If <i>arg</i> is a string, it returns the bank of the section that has that name. If <i>arg</i> is a label, it returns the bank number the label is in. The result may be constant if rgbasm is able to compute it.
SECTION (<i>symbol</i>)	Returns the name of the section that <i>symbol</i> is defined in. <i>symbol</i> must have been defined already. If <i>arg</i> is the symbol @, this function returns the name of the current section.
SIZEOF (<i>arg</i>)	If <i>arg</i> is a string, this function returns the size of the section named <i>arg</i> . The result is a known constant if the section is already "completed" (i.e. not a UNION or FRAGMENT and no longer the current section). If <i>arg</i> is a section type keyword, it returns the size of that section type. The result is not constant, since only RGLINK can compute its value.
STARTOF (<i>arg</i>)	If <i>arg</i> is an 8-bit or 16-bit register, it returns the size of that register. If <i>arg</i> is a string, this function returns the starting address of the section named <i>arg</i> . The result is a known constant if the section has a fixed address. If <i>arg</i> is a section type keyword, it returns the starting address of that section type. The result is not constant, since only RGLINK can compute its value.

SECTIONS

Before you can start writing code, you must define a section. This tells the assembler what kind of information follows and where to put it.

```
SECTION name, type
SECTION name, type, options
SECTION name, type[addr]
SECTION name, type[addr], options
```

name is a string enclosed in double quotes, which is the name of the section. If the type doesn't match, an error occurs. Each section must have a unique name, even across different source files, or the linker will treat it as an error.

Possible section *types* are as follows:

- ROM0** A ROM section. *addr* can range from \$0000 to \$3FFF, or \$0000 to \$7FFF if tiny ROM mode is enabled in the linker.
- ROMX** A banked ROM section. *addr* can range from \$4000 to \$7FFF. Becomes an alias for **ROM0** if tiny ROM mode is enabled in the linker.
- VRAM** A banked video RAM section. *addr* can range from \$8000 to \$9FFF. *bank* can be 0 or 1, but bank 1 is unavailable if DMG mode is enabled in the linker.
- SRAM** A banked external (save) RAM section. *addr* can range from \$A000 to \$BFFF.
- WRAM0** A general-purpose RAM section. *addr* can range from \$C000 to \$CFFF, or \$C000 to \$DFFF if WRAM0 mode is enabled in the linker.
- WRAMX** A banked general-purpose RAM section. *addr* can range from \$D000 to \$DFFF. *bank* can range from 1 to 7. Becomes an alias for **WRAM0** if WRAM0 mode is enabled in the linker.
- OAM** An object attribute RAM section. *addr* can range from \$FE00 to \$FE9F.
- HRAM** A high RAM section. *addr* can range from \$FF80 to \$FFFE.

RGBDS produces ROMs, which means that code and data can only be placed in **ROM0** and **ROMX** sections. The other RAM section types are for statically allocated labels. If you need code or data in RAM, you will need to copy it from ROM to RAM yourself. See “RAM code” for an example of how to conveniently do that with a **LOAD** block.

options are comma-separated and may include:

BANK[*bank*]

Specify which *bank* for the linker to place the section in. See above for possible values for *bank*, depending on *type*.

ALIGN[*align*, *offset*]

Place the section at an address whose *align* least-significant bits are equal to *offset*. Note that **ALIGN**[*align*] is a shorthand for **ALIGN**[*align*, 0]. This option can be used with [*addr*], as long as they don't contradict each other. It's also possible to request alignment in the middle of a section; see “Requesting alignment” below.

If [*addr*] is not specified, the section is considered “floating”; the linker will automatically calculate an appropriate address for the section. Similarly, if **BANK**[*bank*] is not specified, the linker will automatically find a bank with enough space.

Sections can also be placed by using a linker script file. The format is described in *rgblink*(5). They allow the user to place floating sections in the desired bank in the order specified in the script. This is useful if the sections can't be placed at an address manually because the size may change, but they have to be together.

Section examples:

```
SECTION "Cool Stuff", ROMX
```

This switches to the section called “CoolStuff”, creating it if it doesn't already exist. It can end up in any ROM bank. Code and data may follow.

If it is needed, the base address of the section can be specified:

```
SECTION "Cool Stuff", ROMX[$4567]
```

An example with a fixed bank:

```
SECTION "Cool Stuff", ROMX[$4567], BANK[3]
```

And if you want to force only the section's bank, and not its position within the bank, that's also possible:

```
SECTION "Cool Stuff", ROMX, BANK[7]
```

Alignment examples: The first one could be useful for defining an OAM buffer to be DMA'd, since it must be aligned to 256 bytes. The second could also be appropriate for GBC HDMA, or for an optimized copy code that requires alignment.

```
SECTION "OAM Data", WRAM0, ALIGN[8] ; align to 256 bytes
SECTION "VRAM Data", ROMX, BANK[2], ALIGN[4] ; align to 16 bytes
```

The current section can be ended without starting a new section by using **ENDSECTION**. This directive will clear the section context, so you can no longer write code until you start another section. It can be useful to avoid accidentally defining code or data in the wrong section.

Section stack

POPS and **PUSHS** provide the interface to the section stack. The number of entries in the stack is limited only by the amount of memory in your machine.

PUSHS will push the current section context on the section stack. **POPS** can then later be used to restore it. Useful for defining sections in included files when you don't want to override the section context at the point the file was included.

PUSHS can also take the same arguments as **SECTION**, in order to push the current section context and define a new section at the same time:

```
SECTION "Code", ROM0
Function:
    ld a, 42
    PUSHS "Variables", WRAM0
        wAnswer: db
    POPS
    ld [wAnswer], a
```

RAM code

Sometimes you want to have some code (or data) in RAM, e.g. for self-modifying code. But you can't just put it directly in a RAM section; you have to store it in ROM and copy it to RAM at some point. This means that the code will be executed at a different address range than where it's defined, which can be inconvenient for references to labels within that code. This situation is what **LOAD** blocks are designed for. Here's an example of how to use them:

```
SECTION "LOAD example", ROMX
CopyCode:
    ld de, RAMCode
    ld hl, RAMLocation
    ld c, RAMCode.end - RAMCode
.loop
    ld a, [de]
    inc de
    ld [hli], a
    dec c
    jr nz, .loop
    ret

RAMCode:
    LOAD "RAM code", WRAM0
RAMLocation:
```

```

        ld hl, .string
        ld de, $9864
    .copy
        ld a, [hli]
        ld [de], a
        inc de
        and a
        jr nz, .copy
        ret

    .string
        db "Hello World!\0"
    ENDL
    .end

```

A **LOAD** block feels similar to a **SECTION** declaration because it creates a new one. All data and code generated within such a block is placed in the current section like usual, but all labels are created as if they were placed in this newly-created section.

In the example above, all of the code and data will end up in the “LOAD example” section. You will notice the ‘RAMCode’ and ‘RAMLocation’ labels. The former is situated in ROM, where the code is stored, the latter in RAM, where the code will be loaded.

You cannot nest **LOAD** blocks, nor can you change or stop the current section within them.

The current **LOAD** block can be ended by using **ENDL**. This directive is only necessary if you want to resume writing code in its containing ROM section. Any of **LOAD**, **SECTION**, **ENDSECTION**, or **POPS** will end the current **LOAD** block before performing its own function.

LOAD blocks can use the **UNION** or **FRAGMENT** modifiers as described in “Unionized sections” below.

Unionized sections

When you’re tight on RAM, you may want to define overlapping static memory allocations, as explained in the “Allocating overlapping spaces in RAM” section. However, a **UNION** only works within a single file, so it can’t be used e.g. to define temporary variables across several files, all of which use the same statically allocated memory. Unionized sections solve this problem. To declare a unionized section, add a **UNION** keyword after the **SECTION** one; the declaration is otherwise not different. Unionized sections follow some different rules from normal sections:

- The same unionized section (i.e. having the same name) can be declared several times per **rgbasm** invocation, and across several invocations. Different declarations are treated and merged identically whether within the same invocation, or different ones.
- If one section has been declared as unionized, all sections with the same name must be declared unionized as well.
- All declarations must have the same type. For example, even if *rgbblink(1)*’s **-w** flag is used, **WRAM0** and **WRAMX** types are still considered different.
- Different constraints (alignment, bank, etc.) can be specified for each unionized section declaration, but they must all be compatible. For example, alignment must be compatible with any fixed address, all specified banks must be the same, etc.
- Unionized sections cannot have type **ROM0** or **ROMX**.

Different declarations of the same unionized section are not appended, but instead overlaid on top of each other, just like “Allocating overlapping spaces in RAM”. Similarly, the size of an unionized section is the largest of all its declarations.

Section fragments

Section fragments are sections with a small twist: when several fragments with the same name are encountered, they are concatenated into one section instead of producing an error, even across multiple object files. This works within the same file (paralleling the behavior "plain" sections has in previous versions), but also across object files. To declare a section fragment, add a **FRAGMENT** keyword after the **SECTION** one; the declaration is otherwise not different. However, similarly to "Unionized sections", some rules must be followed:

- If one section has been declared as fragment, all sections with the same name must be declared fragments as well.
- All declarations must have the same type. For example, even if *rgblink(1)*'s **-w** flag is used, **WRAM0** and **WRAMX** types are still considered different.
- Different constraints (alignment, bank, etc.) can be specified for each section fragment declaration, but they must all be compatible. For example, alignment must be compatible with any fixed address, all specified banks must be the same, etc.
- A section fragment may not be unionized; after all, that wouldn't make much sense.

When RGBASM merges two fragments, the one encountered later is appended to the one encountered earlier.

When RGBLINK merges two fragments, the one whose file was specified last is appended to the one whose file was specified first. For example, assuming *bar.o*, *baz.o*, and *foo.o* all contain a fragment with the same name, the command

```
rgblink -o rom.gb baz.o foo.o bar.o
```

would produce the fragment from *baz.o* first, followed by the one from *foo.o*, and the one from *bar.o* last.

Fragment literals

Fragment literals are useful for short blocks of code or data that are only referenced once. They are section fragments created by surrounding instructions or directives with '[[' double brackets ']]', without a separate **SECTION FRAGMENT** declaration.

The content of a fragment literal becomes a **SECTION FRAGMENT**, sharing the same name and bank as its parent ROM section, but without any other constraints. The parent section also becomes a **FRAGMENT** if it was not one already, so that it can be merged with its fragment literals. RGBLINK merges the fragments in no particular order.

A fragment literal can take the place of any 16-bit integer constant *n16* from the *gbz80(7)* documentation, as well as a **DW** item. The fragment literal then evaluates to its starting address. For example, you can **CALL** or **JP** to a fragment literal.

This code using named labels:

```
DataTable:
    dw First
    dw Second
    dw Third
First: db 1
Second: db 4
Third: db 9
Routine:
    push hl
    ld hl, Left
    jr z, .got_it
    ld hl, Right
.got_it
    call .print
    pop hl
```

```

    ret
.print:
    ld de, $1003
    ld bc, STARTOF(VRAM)
    jp Print
Left:  db "left\0"
Right: db "right\0"

```

is equivalent to this code using fragment literals:

```

DataTable:
    dw [[ db 1 ]]
    dw [[ db 4 ]]
    dw [[ db 9 ]]
Routine:
    push hl
    ld hl, [[ db "left\0" ]]
    jr z, .got_it
    ld hl, [[ db "right\0" ]]
.got_it
    call [[
        ld de, $1003
        ld bc, STARTOF(VRAM)
        jp Print
    ]]
    pop hl
    ret

```

The difference is that the example using fragment literals does not declare a particular order for its pieces.

Fragment literals can be arbitrarily nested, so extreme use cases are *technically* possible. This code using named labels:

```

    dw FortyTwo
FortyTwo:
    call Sub1
    jr Sub2
Sub1:
    ld a, [Twenty]
    ret
Twenty: db 20
Sub2:
    jp Sub3
Sub3:
    call Sub1
    inc a
    add a
    ret

```

is equivalent to this code using fragment literals:

```

dw [[
    call [[
        Sub1: ld a, [ [[db 20]] ] :: ret
    ]]
    jr [[
        jp [[ call Sub1 :: inc a :: add a :: ret ]]
    ]]
]]

```

]]

SYMBOLS

RGBDS supports several types of symbols:

Label Numeric symbol designating a memory location. May or may not have a value known at assembly time.

Constant Numeric symbol whose value has to be known at assembly time.

Macro A block of **rgbasm** code that can be invoked later.

String A text string that can be expanded later, similarly to a macro.

Symbol names can contain ASCII letters, numbers, underscores ‘_’, hashes ‘#’, dollar signs ‘\$’, and at signs ‘@’. However, they must begin with either a letter or an underscore. Additionally, label names can contain up to a single dot ‘.’, which may not be the first character.

A symbol cannot have the same name as a reserved keyword, unless its name is a “raw identifier” prefixed by a hash ‘#’. For example, **#load** denotes a symbol named **load**, and **#LOAD** denotes a different symbol named **LOAD**; in both cases the ‘#’ prevents them from being treated as the keyword **LOAD**.

Labels

One of the assembler’s main tasks is to keep track of addresses for you, so you can work with meaningful names instead of “magic” numbers. Labels enable just that: a label ties a name to a specific location within a section. A label resolves to a bank and address, determined at the same time as its parent section’s (see further in this section).

A label is defined by writing its name at the beginning of a line, followed by one or two colons, without any whitespace between the label name and the colon(s). Declaring a label (global or local) with two colons ‘::’ will define and **EXPORT** it at the same time. (See “Exporting and importing symbols” below). When defining a local label, the colon can be omitted, and **rgbasm** will act as if there was only one.

A label is said to be *local* if its name contains a dot ‘.’; otherwise, it is said to be *global* (not to be mistaken with “exported”, explained in “Exporting and importing symbols” below). More than one dot in label names is not allowed.

For convenience, local labels can use a shorthand syntax: when a symbol name starting with a dot is found (for example, inside an expression, or when declaring a label), then the current “label scope” is implicitly prepended.

Defining a global label sets it as the current “label scope”, until the next global label definition, or the end of the current section.

Here are some examples of label definitions:

```
GlobalLabel:
AnotherGlobal:
.locallabel ; This defines "AnotherGlobal.locallabel"
.another_local:
AnotherGlobal.with_another_local:
ThisWillBeExported:: ; Note the two colons
ThisWillBeExported.too::
```

In a numeric expression, a label evaluates to its address in memory. (To obtain its bank, use the **BANK()** function described in “Other functions”). For example, given the following, **ld de, vPlayerTiles** would be equivalent to **ld de, \$80C0** assuming the section ends up at **\$80C0**:

```
SECTION "Player tiles", VRAM
vPlayerTiles:
    ds 6 * 16
.end
```

A label's location (and thus value) is usually not determined until the linking stage, so labels usually cannot be used as constants. However, if the section in which the label is defined has a fixed base address, its value is known at assembly time.

Also, while **rgbasm** obviously can compute the difference between two labels if both are constant, it is also able to compute the difference between two non-constant labels if they both belong to the same section, such as `PlayerTiles` and `PlayerTiles.end` above.

Anonymous labels

Anonymous labels are useful for short blocks of code. They are defined like normal labels, but without a name before the colon. Anonymous labels are independent of label scoping, so defining one does not change the scoped label, and referencing one is not affected by the current scoped label.

Anonymous labels are referenced using a colon ':' followed by pluses '+' or minuses '-'. Thus `:+` references the next one after the expression, `:++` the one after that; `:-` references the one before the expression; and so on.

```

        ld hl, :++
:      ld a, [hli] ; referenced by "jr nz"
        ldh [c], a
        dec c
        jr nz, :-
        ret

:      ; referenced by "ld hl"
        dw $7FFF, $1061, $03E0, $58A5

```

Variables

An equal sign '=' is used to define mutable numeric symbols. Unlike the other symbols described below, variables can be redefined. This is useful for internal symbols in macros, for counters, etc.

```

DEF ARRAY_SIZE EQU 4
DEF COUNT = 2
DEF COUNT = 3
DEF COUNT = ARRAY_SIZE + COUNT
DEF COUNT *= 2
; COUNT now has the value 14

```

Note that colons ':' following the name are not allowed.

Variables can be conveniently redefined by compound assignment operators like in C:

Operator	Meaning
<code>+=</code> <code>-=</code>	Compound plus/minus
<code>*=</code> <code>/=</code> <code>%=</code>	Compound multiply/divide/modulo
<code><<=</code> <code>>>=</code>	Compound shift left/right
<code>&=</code> <code> =</code> <code>^=</code>	Compound and/or/xor

Examples:

```

DEF x = 10
DEF x += 1 ; x == 11
DEF y = x - 1 ; y == 10
DEF y *= 2 ; y == 20
DEF y >>= 1 ; y == 10
DEF x ^= y ; x == 1

```

Declaring a variable with **EXPORT DEF** or **EXPORT REDEF** will define and **EXPORT** it at the same time. (See "Exporting and importing symbols" below).

Numeric constants

EQU is used to define numeric constant symbols. Unlike '=' above, constants defined this way cannot be redefined. These constants can be used for unchanging values such as properties of the hardware.

```
def SCREEN_WIDTH equ 160 ; In pixels
def SCREEN_HEIGHT equ 144
```

Note that colons ':' following the name are not allowed.

If you *really* need to, the **REDEF** keyword will define or redefine a numeric constant symbol. (It can also be used for variables, although it's not necessary since they are mutable.) This can be used, for example, to update a constant using a macro, without making it mutable in general.

```
def NUM_ITEMS equ 0
MACRO add_item
    redef NUM_ITEMS equ NUM_ITEMS + 1
    def ITEM_{02x:NUM_ITEMS} equ \1
ENDM
    add_item 1
    add_item 4
    add_item 9
    add_item 16
    assert NUM_ITEMS == 4
    assert ITEM_04 == 16
```

Declaring a numeric constant with **EXPORT DEF** or **EXPORT REDEF** will define and **EXPORT** it at the same time. (See "Exporting and importing symbols" below).

Offset constants

The RS group of commands is a handy way of defining structure offsets:

```
RSRESET
DEF str_pStuff RW 1
DEF str_tData RB 256
DEF str_bCount RB 1
DEF str_SIZEOF RB 0
```

The example defines four constants as if by:

```
DEF str_pStuff EQU 0
DEF str_tData EQU 2
DEF str_bCount EQU 258
DEF str_SIZEOF EQU 259
```

There are five commands in the RS group of commands:

Command	Meaning
RSRESET	Equivalent to RSSET 0 .
RSSET <i>constexpr</i>	Sets the _RS counter to <i>constexpr</i> .
DEF <i>name</i> RB <i>constexpr</i>	Sets <i>name</i> to _RS and then adds <i>constexpr</i> to _RS .
DEF <i>name</i> RW <i>constexpr</i>	Sets <i>name</i> to _RS and then adds <i>constexpr</i> * 2 to _RS .
DEF <i>name</i> RL <i>constexpr</i>	Sets <i>name</i> to _RS and then adds <i>constexpr</i> * 4 to _RS .

If the *constexpr* argument to **RB**, **RW**, or **RL** is omitted, it's assumed to be 1.

Note that colons ':' following the name are not allowed.

Declaring an offset constant with **EXPORT DEF** will define and **EXPORT** it at the same time. (See "Exporting and importing symbols" below).

String constants

EQUS is used to define string constant symbols. Wherever the assembler reads a string constant, it gets *expanded*: the symbol's name is replaced with its contents, similarly to **#define** in the C programming language. This expansion is disabled in a few contexts: **DEF**(name), **DEF** name **EQU**/=**EQU**S/etc . . . , **REDEF** name **EQU**/=**EQU**S/etc . . . , **FOR** name, . . . , **PURGE** name, and **MACRO** name will not expand string constants in their names. Expansion is also disabled if the string constant's name is a raw identifier prefixed by a hash '#'.

```
DEF COUNTREG EQU "[hl+]"
    ld a, COUNTREG

DEF PLAYER_NAME EQU "\"John\""
    db PLAYER_NAME
```

This will be interpreted as:

```
ld a, [hl+]
db "John"
```

String constants can also be used to define small one-line macros:

```
DEF pusha EQU "push af\npush bc\npush de\npush hl\n"
```

Note that colons ':' following the name are not allowed.

String constants, like numeric constants, cannot be redefined. However, the **REDEF** keyword will define or redefine a string constant symbol. For example:

```
DEF s EQU "Hello, "
REDEF s EQU "{s}world!"
; prints "Hello, world!"
PRINTLN "{s}\n"
```

String constants can't be exported or imported.

Important note: When a string constant is expanded, its expansion may contain another string constant, which will be expanded as well, and may be recursive. If this creates an infinite loop, **rgbasm** will error out once a certain depth is reached (see the **-r** command-line option in *rgbasm*(1)). The same problem can occur if the expansion of a string constant invokes a macro, which itself expands.

Macros

One of the best features of an assembler is the ability to write macros for it. Macros can be called with arguments, and can react depending on input using **IF** constructs.

```
MACRO my_macro
    ld a, 80
    call MyFunc
ENDM
```

The example above defines **my_macro** as a new macro. String constants are not expanded within the name of the macro.

Macros can't be exported or imported.

Nesting macro definitions is not possible, so this won't work:

```
MACRO outer
    MACRO inner
        PRINTLN "Hello!"
    ENDM ; this actually ends the 'outer' macro...
ENDM    ; ...and then this is a syntax error!
```

But you can work around this limitation using **EQU**, so this will work:

```
MACRO outer
    DEF definition EQU "MACRO inner\nPRINTLN \"Hello!\"\nENDM"
    definition
    PURGE definition
ENDM
```

More about how to define and invoke macros is described in “THE MACRO LANGUAGE” below.

Exporting and importing symbols

Importing and exporting of symbols is a feature that is very useful when your project spans many source files and, for example, you need to jump to a routine defined in another file.

Exporting of symbols has to be done manually, importing is done automatically if **rgbasm** finds a symbol it does not know about.

The following will cause *symbol1*, *symbol2* and so on to be accessible to other files during the link process:

```
EXPORT symbol1 [, symbol2, ...]
```

For example, if you have the following three files:

a.asm:

```
SECTION "a", WRAM0
LabelA:
```

b.asm:

```
SECTION "b", WRAM0
ExportedLabelB1::
ExportedLabelB2:
    EXPORT ExportedLabelB2
```

c.asm:

```
SECTION "c", ROM0[0]
    dw LabelA
    dw ExportedLabelB1
    dw ExportedLabelB2
```

Then c.asm can use ExportedLabelB1 and ExportedLabelB2, but not LabelA, so linking them together will fail:

```
$ rgbasm -o a.o a.asm
$ rgbasm -o b.o b.asm
$ rgbasm -o c.o c.asm
$ rgbasm -o a.o b.o c.o
error: Undefined symbol "LabelA"
    at c.asm(2)
Linking failed with 1 error
```

Purging symbols

PURGE allows you to completely remove a symbol from the symbol table, as if it had never been defined.

```
DEF value EQU 42
PURGE value
DEF value EQU "I'm a string now"
ASSERT DEF(value)
PURGE value
ASSERT !DEF(value)
```

Be very careful when purging symbols that have been referenced in section data, or that have been exported, because it could result in unpredictable errors if something depends on the missing symbol (for example, expressions the linker needs to calculate). Purging labels at all is *not* recommended.

String constants are not expanded within the symbol names.

Predeclared symbols

The following symbols are defined by the assembler:

Name	Type	Contents
@	EQU	PC value (essentially, the current memory address)
.	EQU	The current global label scope
..	EQU	The current local label scope
__SCOPE__	EQU	The innermost current label scope level (empty,., or ..)
_RS	=	_RS Counter
_NARG	EQU	Number of arguments passed to macro, updated by SHIFT
__ISO_8601_LOCAL__	EQU	ISO 8601 timestamp (local)
__ISO_8601_UTC__	EQU	ISO 8601 timestamp (UTC)
__UTC_YEAR__	EQU	Today's year
__UTC_MONTH__	EQU	Today's month number, 1–12
__UTC_DAY__	EQU	Today's day of the month, 1–31
__UTC_HOUR__	EQU	Current hour, 0–23
__UTC_MINUTE__	EQU	Current minute, 0–59
__UTC_SECOND__	EQU	Current second, 0–59
__RGBDS_MAJOR__	EQU	Major version number of RGBDS
__RGBDS_MINOR__	EQU	Minor version number of RGBDS
__RGBDS_PATCH__	EQU	Patch version number of RGBDS
__RGBDS_RC__	EQU	Release candidate ID of RGBDS, not defined for final releases
__RGBDS_VERSION__	EQU	Version of RGBDS, as printed by <code>rgbasm --version</code>

The current time values will be taken from the `SOURCE_DATE_EPOCH` environment variable if that is defined as a UNIX timestamp. Refer to the spec at [reproducible-builds.org: https://reproducible-builds.org/docs/source-date-epoch/](https://reproducible-builds.org/docs/source-date-epoch/).

DEFINING DATA

Defining constant data in ROM

DB defines a list of bytes that will be stored in the final image. Ideal for tables and text.

```
DB 1,2,3,4,"This is a string"
```

Alternatively, you can use **DW** to store a list of words (16-bit) or **DL** to store a list of double-words/longs (32-bit). Both of these write their data in little-endian byte order; for example, `dw $CAFE` is equivalent to `db $FE, $CA` and not `db $CA, $FE`.

Strings are handled a little specially: they first undergo charmap conversion (see “Character maps”), then each resulting character is output individually. For example, under the default charmap, the following two lines are identical:

```
DW "Hello!"
DW "H", "e", "l", "l", "o", "!"
```

If you do not want this special handling, enclose the string in parentheses.

DS can also be used to fill a region of memory with some repeated values. For example:

```
; outputs 3 bytes: $AA, $AA, $AA
DS 3, $AA
; outputs 7 bytes: $BB, $CC, $BB, $CC, $BB, $CC, $BB
DS 7, $BB, $CC
```

You can also use **DB**, **DW** and **DL** without arguments. This works exactly like **DS 1**, **DS 2** and **DS 4** respectively. Consequently, no-argument **DB**, **DW** and **DL** can be used in a **WRAM0** / **WRAMX** / **HRAM** / **VRAM** / **SRAM** section.

Including binary data files

You probably have some graphics, level data, etc. you'd like to include. Use **INCBIN** to include a raw binary file as it is. If the file isn't found in the current directory, the include-path list passed to *rgbasm*(1)'s **-I** option on the command line will be searched.

```
INCBIN "titlepic.bin"
INCBIN "sprites/hero.bin"
```

You can also include only part of a file with **INCBIN**. The example below includes 256 bytes from *data.bin*, starting from byte 78.

```
INCBIN "data.bin", 78, 256
```

The length argument is optional. If only the start position is specified, the bytes from the start position until the end of the file will be included.

Statically allocating space in RAM

DS statically allocates a number of empty bytes. This is the preferred method of allocating space in a RAM section. You can also use **DB**, **DW** and **DL** without any arguments instead (see “Defining constant data in ROM” below).

```
DS 42 ; Allocates 42 bytes
```

Empty space in RAM sections will not be initialized. In ROM sections, it will be filled with the value passed to the **-p** command-line option, except when using overlays with **-O**.

Instead of an exact number of bytes, you can specify **ALIGN**[*align*, *offset*] to allocate however many bytes are required to align the subsequent data. Thus, '**DS ALIGN**[*align*, *offset*], ...' is equivalent to '**DS** *n*, ...' followed by '**ALIGN**[*align*, *offset*]', where *n* is the minimum value needed to satisfy the **ALIGN** constraint (see “Requesting alignment” below). Note that **ALIGN**[*align*] is a shorthand for **ALIGN**[*align*, 0].

Allocating overlapping spaces in RAM

Unions allow multiple static memory allocations to overlap, like unions in C. This does not increase the amount of memory available, but allows re-using the same memory region for different purposes.

A union starts with a **UNION** keyword, and ends at the corresponding **ENDU** keyword. **NEXTU** separates each block of allocations, and you may use it as many times within a union as necessary.

```
; Let's say PC == $C0DE here
UNION
    ; Here, PC == $C0DE
    wName:: ds 10
    ; Now, PC == $C0E8
    wNickname:: ds 10
    ; PC == $C0F2
NEXTU
    ; PC is back to $C0DE
    wHealth:: dw
    ; PC == $C0E0
    wLives:: db
    ; PC == $C0E1
    ds 7
    ; PC == $C0E8
    wBonus:: db
    ; PC == $C0E9
NEXTU
    ; PC is back to $C0DE again
    wVideoBuffer: ds 16
    ; PC == $C0EE
ENDU
```

```
; Afterward, PC == $C0F2
```

In the example above, ‘wName, wHealth’, and ‘wVideoBuffer’ all have the same value; so do ‘wNickname’ and ‘wBonus’. Thus, keep in mind that `ld [wHealth], a` assembles to the exact same instruction as `ld [wName], a`.

This whole union’s total size is 20 bytes, the size of the largest block (the first one, containing ‘wName’ and ‘wNickname’).

Unions may be nested, with each inner union’s size being determined as above, and affecting its outer union like any other allocation.

Unions may be used in any section, but they may only contain space-allocating directives like **DS** (see “Statically allocating space in RAM”).

Requesting alignment

While **ALIGN** as presented in “SECTIONS” is often useful as-is, sometimes you instead want a particular piece of data (or code) in the middle of the section to be aligned. This is made easier through the use of mid-section **ALIGN** *align, offset*. It will retroactively alter the section’s attributes to ensure that the location the **ALIGN** directive is at, has its *align* lower bits equal to *offset*.

If the constraint cannot be met (for example because the section is fixed at an incompatible address), an error is produced. Note that **ALIGN** *align* is a shorthand for **ALIGN** *align, 0*.

There may be times when you don’t just want to specify an alignment constraint at the current location, but also skip ahead until the constraint can be satisfied. In that case, you can use **DS** **ALIGN**[*align, offset*] to allocate however many bytes are required to align the subsequent data.

If the constraint cannot be met by skipping any amount of space, an error is produced. Note that **ALIGN**[*align*] is a shorthand for **ALIGN**[*align, 0*].

THE MACRO LANGUAGE

Invoking macros

A macro is invoked by using its name at the beginning of a line, like a directive, followed by any comma-separated arguments.

```
add a, b
ld sp, hl
my_macro      ; This will be expanded
sub a, 87
my_macro 42    ; So will this
ret c
my_macro 1, 2 ; And this
```

After **rgbasm** has read the macro invocation line, it will expand the body of the macro (the lines between **MACRO** and **ENDM**) in its place.

Important note: When a macro body is expanded, its expansion may contain another macro invocation, which will be expanded as well, and may be recursive. If this creates an infinite loop, **rgbasm** will error out once a certain depth is reached (see the `-r` command-line option in *rgbasm*(1)). The same problem can occur if the expansion of a macro then expands a string constant, which itself expands.

It’s possible to pass arguments to macros as well!

```
MACRO lb
    ld \1, (\2) << 8 | (\3)
ENDM
lb hl, 20, 18      ; Expands to "ld hl, (20) << 8 | (18)"
lb de, 3 + 1, NUM**2 ; Expands to "ld de, (3 + 1) << 8 | (NUM**2)"
```

You expand the arguments inside the macro body by using the escape sequences `\1` through `\9`, `\1` being the first argument, `\2` being the second, and so on. Since there are only nine digits, you can only use the first nine macro arguments that way. To use the rest, you put the argument number in angle brackets, like

`\<10>`.

This bracketed syntax supports decimal numbers and numeric symbols, where negative values count from the last argument. For example, `\<_NARG>` or `\<-1>` will get the last argument.

Other macro arguments and symbol interpolations will also be expanded inside the angle brackets. For example, if `'\1'` is `'13'`, then `\<\1>` inside the macro body will expand to `\<13>`. Or if `DEF v10 = 42` and `DEF x = 10`, then `\<v{d:x}>` will expand to `\<42>`.

Macro arguments are passed as string constants, although there's no need to enclose them in quotes. Thus, arguments are not evaluated as expressions, but instead are expanded directly inside the macro body. This means that they support all the escape sequences of strings (see "String expressions" above), as well as some of their own:

Sequence	Meaning
<code>'\,</code>	Comma (does not terminate the argument)'
<code>'\('\)</code>	Open parenthesis (does not start enclosing argument contents) Close parenthesis (does not end enclosing argument contents)'

Line continuations work as usual inside macros or lists of macro arguments. However, some characters need to be escaped, as in the following example:

```
MACRO PrintMacro1
    PRINTLN STRCAT(\1)
ENDM
    PrintMacro1 "Hello ", \
                "world"

MACRO PrintMacro2
    PRINT \1
ENDM
    PrintMacro2 STRCAT("Hello ", \
                        "world\n")
```

The comma in `PrintMacro1` needs to be escaped to prevent it from starting another macro argument. The comma in `PrintMacro2` does not need escaping because it is inside parentheses, similar to macro arguments in the C programming language. The backslash in `'\n'` also does not need escaping because quoted string literals work as usual inside macro arguments.

Since macro arguments are expanded directly, it's often a good idea to put parentheses around them if they're meant as part of a numeric expression. For instance, consider the following:

```
MACRO print_double
    PRINTLN \1 * 3
ENDM
    print_double 1 + 2
```

The body will expand to `PRINTLN 1 + 2 * 3`, which will print 7 and not 9 as you might have expected.

The **SHIFT** directive is only available inside macro bodies. It shifts the argument numbers by one to the left, so what was `\2` is now `\1`, what was `\3` is now `\2`, and so forth. (What was `\1` is no longer accessible, so `_NARG` is decreased by 1.)

SHIFT can also take an integer parameter to shift that many times instead of once. A negative parameter will shift the arguments to the right, which can regain access to previously shifted ones.

SHIFT is especially useful in **REPT** loops to iterate over different arguments, evaluating the same loop body each time.

There are some escape sequences which are only valid inside the body of a macro:

Sequence	Meaning
'1' – '9'	The 1st–9th macro argument
\<...>	Further macro arguments
'\#'	All <code>_NARG</code> macro arguments, separated by commas
'\@'	Unique symbol name affix (see below)

The `\@` escape sequence is often useful in macros which define symbols. Suppose your macro expands to a loop of assembly code:

```
MACRO loop_c_times
    xor a, a
    .loop
        ld [hl+], a
        dec c
        jr nz, .loop
    ENDM
```

If you use this macro more than once in the same label scope, it will define `.loop` twice, which is an error. To work around this problem, you can use `\@` as a label suffix:

```
MACRO loop_c_times_fixed
    xor a, a
    .loop\@
        ld [hl+], a
        dec c
        jr nz, .loop\@
    ENDM
```

This will expand to a different value in each invocation, similar to **gensym** in the Lisp programming language.

`\@` also works in **REPT** blocks, expanding to a different value in each iteration.

Automatically repeating blocks of code

Suppose you want to unroll a time-consuming loop without copy-pasting it. **REPT** is here for that purpose. Everything between **REPT** and the matching **ENDR** will be repeated a number of times just as if you had done a copy/paste operation yourself. The following example will assemble `add a, c` four times:

```
REPT 4
    add a, c
ENDR
```

You can also use **REPT** to generate tables on the fly:

```
; Generate a table of square values from 0**2 = 0 to 100**2 = 10000
DEF x = 0
REPT 101
    dw x * x
    DEF x += 1
ENDR
```

As in macros, you can also use the escape sequence `\@`. **REPT** blocks can be nested.

A common pattern is to repeat a block for each value in some range. **FOR** is simpler than **REPT** for that purpose. Everything between **FOR** and the matching **ENDR** will be repeated for each value of a given symbol. String constants are not expanded within the symbol name. For example, this code will produce a table of squared values from 0 to 255:

```
FOR N, 256
    dw N * N
ENDR
```

It acts just as if you had done:

```
DEF N = 0
    dw N * N
DEF N = 1
    dw N * N
DEF N = 2
    dw N * N
; ...
DEF N = 255
    dw N * N
DEF N = 256
```

You can customize the range of **FOR** values, similarly to the `range` function in the Python programming language:

Code	Range
FOR <i>V</i> , <i>stop</i>	<i>V</i> increments from 0 to <i>stop</i>
FOR <i>V</i> , <i>start</i> , <i>stop</i>	<i>V</i> increments from <i>start</i> to <i>stop</i>
FOR <i>V</i> , <i>start</i> , <i>stop</i> , <i>step</i>	<i>V</i> goes from <i>start</i> to <i>stop</i> by <i>step</i>

The **FOR** value will be updated by *step* until it reaches or exceeds *stop*, i.e. it covers the half-open range from *start* (inclusive) to *stop* (exclusive). The variable *V* will be assigned this value at the beginning of each new iteration; any changes made to it within the **FOR** loop's body will be overwritten. So the symbol *V* need not be already defined before any iterations of the **FOR** loop, but it must be a variable ("Variables") if so. For example:

```
FOR V, 4, 25, 5
    PRINT "{d:V} "
    DEF V *= 2
ENDR
    PRINTLN "done {d:V}"
```

This will print:

```
4 9 14 19 24 done 29
```

Just like with **REPT** blocks, you can use the escape sequence `\@` inside of **FOR** blocks, and they can be nested.

You can stop a repeating block with the **BREAK** command. A **BREAK** inside of a **REPT** or **FOR** block will interrupt the current iteration and not repeat any more. It will continue running code after the block's **ENDR**. For example:

```
FOR V, 1, 100
    PRINT "{d:V}"
    IF V == 5
        PRINT " stop! "
        BREAK
    ENDC
    PRINT ", "
ENDR
    PRINTLN "done {d:V}"
```

This will print:

```
1, 2, 3, 4, 5 stop! done 5
```

Conditionally assembling blocks of code

The four commands **IF**, **ELIF**, **ELSE**, and **ENDC** let you have **rgbasm** skip over parts of your code depending on a condition. This is a powerful feature commonly used in macros.

```

IF NUM < 0
    PRINTLN "NUM < 0"
ELIF NUM == 0
    PRINTLN "NUM == 0"
ELSE
    PRINTLN "NUM > 0"
ENDC

```

The **ELIF** (standing for "else if") and **ELSE** blocks are optional. **IF** / **ELIF** / **ELSE** / **ENDC** blocks can be nested.

Note that if an **ELSE** block is found before an **ELIF** block, the **ELIF** block will be ignored. All **ELIF** blocks must go before the **ELSE** block. Also, if there is more than one **ELSE** block, all of them but the first one are ignored.

Including other source files

Use **INCLUDE** to process another assembler file and then return to the current file when done. If the file isn't found in the current directory, the include-path list passed to *rgbasm*(1)'s **-I** option on the command line will be searched. You may nest **INCLUDE** calls infinitely (or until you run out of memory, whichever comes first).

```
INCLUDE "irq.inc"
```

You may also implicitly **INCLUDE** a file before the source file with the **-P** option of *rgbasm*(1).

Printing things during assembly

The **PRINT** and **PRINTLN** commands print text and values to the standard output. Useful for debugging macros, or wherever you may feel the need to tell yourself some important information.

```

PRINT "Hello world!\n"
PRINTLN "Hello world!"
PRINT _NARG, " arguments\n"
PRINTLN "sum: ", 2+3, " product: ", 2*3
PRINTLN STRFMT("E = %f", 2.718)

```

PRINT prints out each of its comma-separated arguments. Numbers are printed as unsigned uppercase hexadecimal with a leading '\$'. For different formats, use **STRFMT**.

PRINTLN prints out each of its comma-separated arguments, if any, followed by a newline ('\n').

Aborting the assembly process

FAIL and **WARN** can be used to print errors and warnings respectively during the assembly process. This is especially useful for macros that get an invalid argument. **FAIL** and **WARN** take a string as the only argument and they will print this string out as a normal error with a line number.

FAIL stops assembling immediately while **WARN** shows the message but continues afterwards.

If you need to ensure some assumption is correct when compiling, you can use **ASSERT** and **STATIC_ASSERT**. Syntax examples are given below:

```

Function:
    xor a
ASSERT LOW(MyByte) == 0
    ld h, HIGH(MyByte)
    ld l, a
    ld a, [hli]
; You can also indent this!
    ASSERT BANK(OtherFunction) == BANK(Function)
    call OtherFunction
; Lowercase also works
    ld hl, FirstByte
    ld a, [hli]

```

```

assert FirstByte + 1 == SecondByte
    ld b, [hl]
    ret
.end
; If you specify one, a message will be printed
STATIC_ASSERT .end - Function < 256, "Function is too large!"

```

First, the difference between **ASSERT** and **STATIC_ASSERT** is that the former is evaluated by RGBASM if it can, otherwise by RGBLINK; but the latter is only ever evaluated by RGBASM. If RGBASM cannot compute the value of the argument to **STATIC_ASSERT**, it will produce an error.

Second, as shown above, a string can be optionally added at the end, to give insight into what the assertion is checking.

Finally, you can add one of **WARN**, **FAIL** or **FATAL** as the first optional argument to either **ASSERT** or **STATIC_ASSERT**. If the assertion fails, **WARN** will cause a simple warning (controlled by *rgbasm(1)* flag *-Wassert*) to be emitted; **FAIL** (the default) will cause a non-fatal error; and **FATAL** immediately aborts.

MISCELLANEOUS

Changing options while assembling

OPT can be used to change some of the options during assembling from within the source, instead of defining them on the command-line. (See *rgbasm(1)*).

OPT takes a comma-separated list of options as its argument:

```

PUSHO
    OPT g.o0X, wdiv      ; acts like command-line '-g.o0X -wdiv'
    OPT -wdiv            ; dashes before the options are optional
    DW '..oo00XX        ; uses the graphics constant characters from OPT g
    PRINTLN $80000000/-1 ; prints a warning about division
POPO
    DW '00112233        ; uses the default graphics constant characters
    PRINTLN $80000000/-1 ; no warning by default

```

OPT can modify the options *b*, *g*, *p*, *Q*, *r*, and *w*.

POPO and **PUSHO** provide the interface to the option stack. **PUSHO** will push the current set of options on the option stack. **POPO** can then later be used to restore them. Useful if you want to change some options in an include file and you don't want to destroy the options set by the program that included your file. The stack's number of entries is limited only by the amount of memory in your machine.

PUSHO can also take a comma-separated list of options, to push the current set and apply the argument set at the same time:

```

PUSHO b.#, g.o0#
    DB %..#####..
    DW '..oo00##
POPO

```

Excluding locations from backtraces

Errors and warnings print *backtraces* showing the location in the source file where the problem occurred, tracing the origin of the problem even through a chain of **REPT**, **FOR**, **MACRO**, and **INCLUDE** locations. Sometimes there are locations you would like to ignore; for example, a common utility macro when you only care about the line where the macro is used, or an **INCLUDE** file that only serves to include other files and is just filler in the backtrace.

In those cases, you can *silence* a location with a question mark '?' after the token: all of the locations created by a 'REPT?', 'FOR?', or 'MACRO?' will not be printed, and any location created by a 'INCLUDE?', or a macro invocation whose name is immediately followed by a '?', will not be printed. For example, if this were assembled as *example.asm*:

```
MACRO lb
    assert -128 <= (\2) && (\2) < 256, "\2 is not a byte"
    assert -128 <= (\3) && (\3) < 256, "\3 is not a byte"
    ld \1, (LOW(\2) << 8) | LOW(\3)
ENDM
SECTION "Code", ROM0
    lb hl, $123, $45
```

This would print an error backtrace:

```
error: Assertion failed: $123 is not a byte
    at example.asm::lb(2)
    <- example.asm(7)
```

But if `MACRO` were changed to `MACRO?`, or `lb hl` were changed to `lb? hl`, then the error backtrace would not mention the location within the ‘`lb`’ macro:

```
error: Assertion failed: $123 is not a byte
    at example.asm(7)
```

SEE ALSO

rgbasm(1), *rgbblink*(1), *rgbblink*(5), *rgbfix*(1), *rgbgfx*(1), *gbz80*(7), *rgbasm-old*(5), *rgbds*(5), *rgbds*(7)

HISTORY

rgbasm(1) was originally written by Carsten Sørensen as part of the ASMotor package, and was later repackaged in RGBDS by Justin Lloyd. It is now maintained by a number of contributors at <https://github.com/gbdev/rgbds>.